

Where Sampling-Grid Dependence Enters Selective State-Space Models

Anonymous Authors¹

Abstract

A sequence model can change its output when only the time grid of the same physical input changes; we ask where this dependence enters a selective state-space model (SSM), separating interval splitting, which keeps the held input path, from new observations, which change it. For one layer, a physical step, exact zero-order hold and coefficients that see only the held value give split invariance; in a stack, later layers receive an internal signal that the refined grid samples differently. On one selective SSM with a physical step (official TIDES code) and four seeds on one synthetic task, we separate a model-coupling change at common times, replicated in every seed; a token-mean readout change under a density change, mostly the weighted-average identity of the readout; and a native-versus-resampled task-risk contrast, which changes sign across seeds. One checkpoint on a public activity-recognition dataset (UCI HAR, nine subjects) shows the same separation: under a pooled splitting change of 0.0290 in class-centered logits, the change at the common time points is 0.0018, an order of magnitude below the readout term over the added tokens; the two are not additive shares. The pre-registered cross-entropy difference between token-mean and time-weighted pooling is +0.016 nats with a subject-level interval that includes zero, so no risk difference is established; two further seeds on the same subjects give +0.050 and +0.044 nats as an unpooled sensitivity check. We propose no new architecture.

1. Introduction

Continuous-time state-space layers are applied to sequences by discretizing a linear ODE over each interval between

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

observations. When the step is the physical elapsed time, the recurrence inherits the semantics of the continuous system, which lets S4, LSSL and S5 change the sampling rate at test time or consume irregular timestamps (Gu et al., 2021a;b; Smith et al., 2022). Selective SSMs such as Mamba compute the step from the current token (Gu & Dao, 2023; Dao & Gu, 2024), so it acts as a content gate; TIDES restores a physical step and moves selectivity to the state matrix (Soydan et al., 2026).

We do not propose another architecture. We ask a measurement question: *when only the time grid of the same physical input path changes, where does the output change enter?* Splitting a held interval into m sub-steps that repeat the held value leaves the held input path unchanged; we call an output change at the common times an *artifact*, in the operational sense of a violation of this pre-declared invariance. Adding observations of the same signal changes the observed values and hence the reconstructed input path, so part of the change is a response to a different input; we make no claim about task-relevant or Bayesian information.

Known and not claimed. The physical Δt , exact zero-order-hold (ZOH) formulas and flow composition are standard (Gu et al., 2021b; Smith et al., 2022); that the released Mamba code uses $\bar{B} = \Delta B$ (“exponential-Euler”) is documented by Lahoti et al. (2026); that a continuous-time reading of selective SSMs needs $\Delta \propto dt$ is shown by Cirone et al. (2024); keeping Δ physical with selectivity on the decay rate is the design of Soydan et al. (2026); a sample mean weighting samples rather than time is elementary.

Contributions. None is an architecture.

1. A protocol that separates splitting from new observations, with an exact path/artifact decomposition and explicit denominators (Section 3).
2. Two different causes, stated with their assumptions: resampling of internal signals between layers, which changes a stack of exact-update layers under splitting although the held input is unchanged, with an $O(\text{mesh})$ upper bound for a scalar linear cascade (Section 4.2), and sample-count readouts, which weight tokens instead of time (Section 4.3); toy checks in Section 6.

3. A fixed-protocol seed repetition on the official TIDES code that separates a model-coupling effect, a pooled-weighting effect and a task-risk effect, with cost; an exact signed decomposition shows that the token-mean readout effect is almost entirely a weighted-average identity (Section 5.1); and one pre-registered evaluation of a checkpoint trained on real sensor windows, with a fixed-checkpoint diagnostic that separates the common-time output change from the readout’s averaging over added time points, and a two-seed sensitivity check (Section 5.2).
4. A check that the TIDES paper equations and the released code pair inputs with different intervals (Section 4.4).

2. Related Work

Discretization in structured SSMs. S4 uses the bilinear transform and reports a test-time resolution change obtained by changing its internal step (Gu et al., 2021a); LSSL handles irregular data by discretizing with a different timescale (Gu et al., 2021b); S4D gives bilinear and ZOH rules (Gu et al., 2022). S5 uses ZOH, $\bar{\Lambda} = e^{\Delta\Lambda}$, $\bar{B} = \Lambda^{-1}(\bar{\Lambda} - I)\bar{B}$, rescales Δ globally for zero-shot 8 kHz evaluation of a 16 kHz model, and supplies a per-step Δ_t for an irregular pendulum task, with ablations that drop the interval or append it as a feature (Smith et al., 2022). S4ND notes that only the relative rate matters (Nguyen et al., 2022); event-camera SSMs evaluate at different frequencies (Zubić et al., 2024). These are LTI layers.

Selective SSMs. Mamba makes Δ, B, C input-dependent and states the ZOH rule (Gu & Dao, 2023); Mamba-2 keeps its discretization (Dao & Gu, 2024). Mamba-3 derives right-endpoint ZOH for time-varying systems, names the released rule exponential-Euler, and proposes an exponential-trapezoidal rule that is second order only if its data-dependent weight is $\frac{1}{2} + O(\Delta_t)$ (Lahoti et al., 2026). Liquid-S4 makes the transition input-dependent (Hasani et al., 2022).

Physical time in selective SSMs. TIDES keeps $\tilde{\Delta} \equiv \Delta$, makes $\text{Re } \Lambda, B$ and C input-dependent, and uses $\bar{B}_k = \Lambda_k^{-1}(\bar{\Lambda}_k - I)B_k$ (Soydan et al., 2026). Its Fading Flash diagnostic rescales a global step for fixed-length sequences and shows that a Mamba surrogate with Δ as an input channel fails outside the training range; its random-drop study removes observations and reports accuracy. In the sections we read it neither refines a fixed held path nor decomposes output changes; it is the closest prior work, and we evaluate its official code, not a re-implementation, treating its architecture as prior art.

Continuous-time models. Neural CDEs are driven by an interpolation of the observations (Kidger et al., 2020); Log-NCDEs build on NCDEs, which are robust to irregular sampling (Walker et al., 2024); Cirone et al. (2024) read selective SSMs as linear CDEs in a limit where δ plays the role of dt ; oscillatory SSMs study implicit and IMEX integration (Rusch & Rus, 2024). We did not find, in these sources, a refinement of the same held path for selective SSMs that separates splitting from new observations and treats readouts and inter-layer resampling as entry points; this is absence of evidence, not proof of novelty.

3. Problem Setup

Model and holds. The toy is a scalar-input selective model with N diagonal modes,

$$\dot{h}_n = g(u)(\lambda_n h_n + B_n(u)u), \quad y = \text{Re} \sum_n C_n(u)h_n, \quad (1)$$

with $\text{Re } \lambda_n < 0$, a gate $g(u) = \text{softplus}(wu + \beta)$ shared across modes, and affine $B(u), C(u)$. A grid $0 = t_0 < \dots < t_K = T$ has steps dt_k and mesh δ . We use the right-endpoint hold $\tilde{u}(t) = u(t_k)$ on $(t_{k-1}, t_k]$, which is the convention of the Mamba recurrence and of the TIDES code (Section 4.4). A discrete variant is $h_k = \bar{A}_k h_{k-1} + \bar{B}_k u_k$ with $\Delta_k = dt_k g(u_k)$ (“ $\Delta = dt g$ ”) or $\Delta_k = \tau g(u_k)$ with a fixed τ (“ $\Delta = \tau g$ ”), and exact ZOH, Euler- B ($\bar{B} = \Delta B$) or bilinear rules; τ is the base step.

Grid changes and decomposition. *Splitting* divides base intervals into m equal sub-steps that all receive the held value; *new observations* give each new sub-time the value of the continuous signal. With D_m the outputs of a variant on grid m and Y_m the continuous-time solution of (1) on the held path of grid m , the change $T_m = D_m - D_1$ at the common base times splits exactly as

$$T_m = \underbrace{(Y_m - Y_1)}_{P_m: \text{path}} + \underbrace{[(D_m - Y_m) - (D_1 - Y_1)]}_{R_m: \text{artifact}}, \quad (2)$$

so $\|T_m\|^2 = \|P_m\|^2 + \|R_m\|^2 + 2\langle P_m, R_m \rangle$. Under splitting $P_m = 0$. Because the cross term can be large and of either sign, we never report a ratio of norms as a share or as a causal contribution.

Readouts. We distinguish sample-count readouts ($\sum_k y_k$, $\frac{1}{K} \sum_k y_k$), time-weighted endpoint sums (right-endpoint $\frac{1}{T} \sum_k y_k dt_k$ and trapezoidal), and the time average $\frac{1}{T} \int_0^T y dt$ of the output over the held path, which the toy computes in closed form for exact-ZOH states (the “exact integral”).

4. Where Grid Dependence Enters: Analysis

4.1. Single layer

Proposition 4.1 (standard). *Fix a base interval $(t_{k-1}, t_k]$ and suppose (i) every sub-step receives the held value u_k ; (ii) every coefficient (g, B, C, λ) is a function of the held value only; (iii) the step is (sub-step length) $\cdot g(u_k)$; and (iv) $\bar{A}, \bar{B}$ are the exact ZOH formulas. Then, from the same state at t_{k-1} , the state and output at t_k are the same for every m ; by induction the outputs at all base times are split-invariant.*

Proof. Each sub-step is the exact flow Φ_s of $\dot{h} = g(u_k)(\lambda h + B(u_k)u_k)$ over its length s , and $\Phi_{s_1} \circ \Phi_{s_2} = \Phi_{s_1+s_2}$. $\square$

The proposition is textbook; its use is that each entry point violates one assumption. $\Delta = \tau g$ violates (iii): m sub-steps apply m times the decay. Euler- B violates (iv), with local error $|\Delta B u| |\varphi_1(\Delta \lambda) - 1| \leq |B u| \Delta^2 |\lambda|/2$, $\varphi_1(z) = (e^z - 1)/z$, since $e^z - 1 - z = z^2 \int_0^1 (1-s)e^{sz} ds$ gives $|\varphi_1(z) - 1| \leq |z|/2$ for $\text{Re } z \leq 0$. (ii) fails whenever a coefficient sees anything besides the held value: a token-indexed convolution (Mamba’s `d_conv=4`), time or position features, state-dependent transitions, or training-mode normalization statistics over time. Bilinear is stable for $\text{Re } z < 0$, but for real $z < -2$ its transition is negative (sign-alternating transients) and tends to -1 instead of $e^z \rightarrow 0$ (inaccurate stiff decay); its held-input fixed point is exact.

4.2. Several layers: internal resampling

In a stack, layer 2 is driven by the continuous output of layer 1, but a layerwise recurrence gives it only endpoint samples held over each step. Splitting the input grid changes which internal samples it sees, without adding any external observation. The minimal instance is

$$\dot{h}_1 = -a h_1 + b u(t), \quad \dot{h}_2 = -c h_2 + d h_1(t). \quad (3)$$

Proposition 4.2 (standard). *Let u be held on a grid with mesh δ , $|u| \leq U$, $a \geq 0$, $c > 0$ and $h_1(0) = h_2(0) = 0$. (a) The exact coupled solution of (3) at base times is split-invariant. (b) The layerwise computation that is exact in layer 1 and holds $h_1(t_k)$ on $(t_{k-1}, t_k]$ in layer 2 has absolute error $|e_k| \leq |b d| U \delta e^{c\delta}/c$ at every grid time, for every T and also for $a = c$.*

Proof in Appendix F.

Remark 4.3 (scope of Proposition 4.2). The argument is elementary and has not been checked by a third party. A nonzero $h_1(0)$ adds $a|h_1(0)|$ to $2|b|U$; a nonzero $h_2(0)$ does not enter. Uniformity in T needs $c > 0$: as $c \rightarrow 0$ the sum is bounded only by t_k and the bound grows linearly in time; $a < 0$ (an unstable first layer) is not covered. The statement

is for a linear time-invariant scalar cascade; it gives no guarantee for input-dependent coefficients, nonlinear maps between layers, or trained deep models, where the Lipschitz and contraction constants it uses are not controlled.

Conditions (ii)–(iv) of Proposition 4.1 concern each layer’s update and can hold in every layer, but condition (i) holds only for the first: the input of layer 2 is the output of layer 1, which is not constant within a base interval, so a stacked recurrence on the observation grid sees a different layer-2 input after splitting. The resulting change is a reconstruction error of an *internal* signal, not a change of the observed input. Proposition 4.2(b) is an upper bound of order δ for the stated scalar cascade; it does not show that the error is nonzero or of first order for every input (for $u \equiv 0$ it is zero), and it is not a theorem about deep selective networks. Two quantities should be kept apart. The *invariance defect*, the change of the base-time outputs under splitting, vanishes if the inter-layer signal is represented exactly or if the stack runs on a fixed grid that does not depend on the observations (for example the training grid after resampling). The *approximation error* of the layerwise computation relative to the continuous coupled system does not vanish on a fixed grid; for the cascade it is bounded as in Proposition 4.2(b) with the mesh of that grid.

4.3. Readouts

The time average of the output over the held path is additive over sub-intervals, so it is split-invariant whenever the state trajectory is (for example under the conditions of Proposition 4.1); endpoint sums converge to it only where the grid is refined; sample-count readouts weight samples, not time, so the sample mean converges to a density-weighted average. If the first half of $[0, T]$ is split by m and the second half is not, the token mean of outputs that merely repeat the base-grid outputs differs from their base-grid mean by $\frac{m-1}{2(m+1)}(A - B)$, where A and B are the mean base-grid outputs over the two halves ($\frac{7}{18}(A - B)$ for $m=8$); this weighted-average identity involves no property of the model.

4.4. Interval conventions in the TIDES paper and code

The TIDES paper holds u_k forward on $[t_k, t_{k+1})$ with the state in y_k not yet containing u_k , whereas the pinned code holds u_k backward on $(t_{k-1}, t_k]$ with the state in y_k containing it; both are exact ZOH solutions of different input reconstructions, and on an irregular grid the difference is not a relabeling (checked with untrained weights: Appendix E). The released code coincides with the right-endpoint convention used throughout this paper, so our splitting protocol repeats the right-endpoint value accordingly; we make no claim that this affects TIDES’ results.

5. Trained Official TIDES: Synthetic Seed Repetition and One Real-Sensor Evaluation

5.1. Synthetic task: fixed-protocol seed repetition

Setup. *Model:* official TIDES at a pinned commit (MIT), 23,397 parameters, two blocks, input-dependent Re Λ , B , C , exact ZOH, causal, no convolution, a linear per-step head; CPU, float32, 2 threads. *Task:* regression of the teacher (1) with fixed parameters, driven by sine-sum inputs, over $T=8$ on a uniform training grid of $K=32$ steps; labels are the teacher output on the continuous path at the 32 grid times (label s.d. 2.35); trajectory identifiers are split into 512/128 train/calibration before any grid is built, with normalization from the training split only. *Training:* Adam, 2000 steps, batch 32; calibration MSE every 100 steps; the lowest-calibration-MSE state is saved, hashed and frozen before any test trajectory is generated. A development run with seed 0 (Appendix C) chose the questions below; the repetition was then fixed: seeds 1–4 for the initialization and minibatch order, train and calibration data hash-identical to the development run, and the repetition code reproduces every tensor of the development checkpoint (maximum absolute difference 0). All seeds share one new test set of 256 trajectories from the same generator, with identifiers disjoint from all earlier ones; no seed was replaced or added.

Conditions, arms and estimands. The held-path refinements are C0 (training grid), S2/S4/S8 (every interval split with the held value repeated) and H8 (only the first half split by 8, a density change on the same held path); J1 (K random observation times) changes the observed path and is reported separately. *Native* evaluation feeds the condition’s grid with physical step sizes; *resampled* evaluation carries the last observation forward onto the training grid using only observed times and values. On S and H conditions resampling reproduces the C0 input exactly, so its zero discrepancy is an identity control, not evidence of superiority. For a sequence-level readout let $\bar{y}_i = T^{-1} \int_0^T y_i dt$ be the pooled label (Simpson’s rule on the teacher’s RK4 nodes) and $\hat{y}_{i,k}(c)$ the native outputs on condition c with steps $dt_k(c)$ and L_c tokens. We compare the token mean $p_i^{\text{mean}}(c) = L_c^{-1} \sum_k \hat{y}_{i,k}(c)$, the pooling of the TIDES classifier head (here applied to the outputs of a linear per-step head, which equals pooling the features first), with the time-weighted sum $p_i^{\text{time}}(c) = T^{-1} \sum_k \hat{y}_{i,k}(c) dt_k(c)$ of the *same* outputs. The latter is a pooling-only ablation, not part of the official model and not an exact integral; the resampled arm pools its K outputs by $p_i^{\text{res}}(c) = K^{-1} \sum_k \hat{y}_{i,k}^{\text{res}}(c)$. On the uniform C0 grid all three equal $p_i(\text{C0})$; the pooled error of an arm is $n^{-1} \sum_i |p_i(c) - \bar{y}_i|$. The primary estimand, chosen after the development run, is the pooled-readout effect $D_s = \frac{1}{n} \sum_i (|p_{s,i}^{\text{mean}}(\text{H8}) - p_{s,i}(\text{C0})| - |p_{s,i}^{\text{time}}(\text{H8}) - p_{s,i}(\text{C0})|)$ in label units. Secondary estimands are the final-output dis-

Table 1. Model-coupling effect (P1-REAL-02, training seeds 1–4; same task, data and checkpoint rule; shared test set of 256 trajectories, ids 10000–10255). Step and calib. MSE: selected checkpoint and its calibration MSE (label units squared). S8 disc.: mean over trajectories of the relative ℓ_2 change of the final per-step outputs at the 32 training-grid times when every held interval is split into 8 sub-steps, against C0 (float32). Brackets: 95% paired bootstrap CI over test trajectories with the seed fixed.

seed	step	calib. MSE	S8 disc. [95% CI]
1	1300	0.0127	0.0113 [0.0106, 0.0121]
2	1600	0.0135	0.0219 [0.0203, 0.0236]
3	800	0.0168	0.0043 [0.0041, 0.0045]
4	1800	0.0125	0.0147 [0.0140, 0.0153]

crepancy under S8, $\|\hat{y}(\text{S8}) - \hat{y}(\text{C0})\|_2 / \|\hat{y}(\text{C0})\|_2$ over the 32 grid times (also with weights and inputs cast to float64); the loss contrast $r_s = (\bar{L}_{\text{native}} - \bar{L}_{\text{resampled}}) / \bar{L}_{\text{resampled}}$ of the per-trajectory MSE; and wall time. Uncertainty is a paired bootstrap over the 256 test trajectories with the seed fixed, with ratios recomputed in every draw. Every interval is pointwise, for one estimand and one seed; it is not adjusted for the number of secondary conditions and seeds, and for a non-negative discrepancy an interval excluding zero shows that the change is nonzero, not what causes it. Because the seeds share the trajectories (a crossed design), we report the seeds separately and give only their range, not a pooled interval. A loss difference is called negligible only if the interval of r_s lies inside $\pm 5\%$, a convention (about 2.5% in RMSE) rather than an application tolerance; non-significance is never read as equivalence.

Model coupling. The S8 discrepancy is 4.3×10^{-3} –0.022 in the four seeds, each interval excluding zero (Table 1); it varies 5.1-fold between seeds (the earliest-selected checkpoint, step 800, has the smallest value, a relation not tested). Casting weights and inputs to float64 changes the mean by at most 5.2×10^{-6} and any per-trajectory value by at most 1.8×10^{-5} , far below the smallest per-trajectory discrepancy 1.6×10^{-3} , so this is not float32 rounding; the pre-registered 10^{-4} threshold came from a scalar emulation (Appendix G) and is not a floor of a trained 256-token network. In the development checkpoint the change is absent from the encoder and first block ($\leq 7.2 \times 10^{-6}$ in float32, 1.8×10^{-14} in float64) and enters in the second block (4.1×10^{-3} – 7.2×10^{-3}), qualitatively consistent with Propositions 4.1 and 4.2; this is exploratory and rests on one checkpoint (Appendix C).

Pooled weighting. D_s is positive with an interval excluding zero in all four seeds (0.578–0.602, smallest lower bound 0.512; Table 2), so the pre-specified replication criterion is met, but the cause is mostly the readout’s weights: per trajectory the token-mean change splits exactly into $W = \frac{7}{18}(A - B)$, computed from the unchanged C0 outputs

Table 2. Readout decomposition under the same-path density change H8 (first half of $[0, T]$ split by 8; 144 tokens), training seeds 1–4, label units, means over the 256 test trajectories. D_s (primary, chosen after the development run): mean of $|p^{\text{mean}}(\text{H8}) - p(\text{C0})| - |p^{\text{time}}(\text{H8}) - p(\text{C0})|$. Per trajectory, exactly, $p^{\text{mean}}(\text{H8}) - p(\text{C0}) = W + M$ and $p^{\text{time}}(\text{H8}) - p(\text{C0}) = Q$ with $W = \sum_j (a_j - w_j)b_j = \frac{7}{18}(A - B)$, $M = \sum_j a_j(z_j - b_j)$, $Q = \sum_j w_j(z_j - b_j)$; b_j is the C0 output of token j ’s base interval, z_j the native H8 output, $a_j = 1/144$, $w_j = dt_j/T$, and A, B are the mean C0 outputs over the first and second half of $[0, T]$. W uses only C0 outputs: it is the same for any model whose H8 outputs repeat the C0 output within each base interval. The signed terms are not an additive decomposition of the absolute-value estimand D_s . Brackets: 95% paired bootstrap CI (seed fixed).

seed	D_s [95% CI]	$p^{\text{mean}} - p(\text{C0})$	W	M [95% CI]	Q [95% CI]	$ Q $
1	0.578 [0.512, 0.649]	+0.462	+0.476	-0.014 [-0.018, -0.010]	-0.012 [-0.015, -0.009]	0.019
2	0.584 [0.517, 0.655]	+0.473	+0.488	-0.016 [-0.019, -0.012]	-0.015 [-0.017, -0.012]	0.023
3	0.602 [0.536, 0.673]	+0.488	+0.491	-0.003 [-0.007, +0.001]	-0.001 [-0.003, +0.001]	0.013
4	0.602 [0.536, 0.672]	+0.497	+0.486	+0.011 [+0.008, +0.015]	+0.014 [+0.012, +0.016]	0.018

Table 3. Task-risk effect: loss contrast $r_s = (\bar{L}_{\text{native}} - \bar{L}_{\text{resampled}})/\bar{L}_{\text{resampled}}$ in percent (per-trajectory MSE at the 32 training-grid times; numerator and denominator recomputed in every bootstrap draw), training seeds 1–4. Verdict: = CI inside $\pm 5\%$; < / > native lower / higher (CI excludes 0); << / >> also beyond the margin; ? inconclusive. The $\pm 5\%$ margin is a convention. S2, S8 and H8 keep the held path; J1 (random observation times) changes the observed path. Timing: Table 7.

seed	$r_s(\text{S2})$ [95% CI]		$r_s(\text{S8})$ [95% CI]		$r_s(\text{H8})$ [95% CI]		$r_s(\text{J1})$ [95% CI]	
1	-4.3 [-5.8, -2.7]	<	-6.2 [-8.8, -3.4]	<	-4.9 [-7.2, -2.3]	<	-21.9 [-32.0, -8.8]	<<
2	-2.1 [-5.3, +1.3]	?	+3.1 [-2.5, +9.4]	?	+6.7 [+1.8, +11.9]	>	-21.2 [-31.7, -6.5]	<<
3	-1.2 [-2.4, -0.2]	=	-1.9 [-3.9, -0.1]	=	-0.6 [-1.7, +0.6]	=	-21.8 [-31.9, -8.9]	<<
4	-5.4 [-8.0, -2.8]	<	-4.5 [-9.6, +0.7]	?	-2.1 [-5.7, +1.9]	?	-22.6 [-32.9, -9.0]	<<

(Section 4.3), and a model term M . The change averages +0.46 to +0.50, W averages +0.48 to +0.49 and M averages -0.016 to +0.011; W has the sign of the change in at least 98% of trajectories, and a token mean of outputs that merely repeat C0 would err by 0.56–0.59 against the pooled label, as much as the native token mean (0.57–0.59). The effect is therefore sample-count pooling of a non-constant output, which any per-token model shows, not a failure specific to selective SSMs. $|M|$ averages 0.023–0.028, and 0.0005–0.0049 on the tokens that end a base interval, where it is the change of the model output itself. Time weighting reduces the mean absolute change against C0 from 0.60–0.62 to 0.013–0.023 but does not remove it, and its pooled error (0.068–0.094) stays close to the resampled arm’s (0.066–0.096). W , M and Q come from one frozen forward pass per checkpoint that reproduces the stored pooled values exactly (Appendix A); the signed terms do not decompose the absolute-value estimand D_s .

Task risk: the accuracy consequence does not replicate.

The development run had suggested slightly better native accuracy on refined grids. Across seeds 1–4, r_s ranges from -6.2% to +3.1% on S8 and from -4.9% to +6.7% on H8 (Table 3): seed 2 is worse on H8 with an interval that excludes zero, seed 3 shows no practical difference on every held-path condition, and the rest are mixed, so we report no stable difference. On J1 native MSE is lower in every seed (r_s from -22.6% to -21.2%), but the development checkpoint gives -20.9% here and -6.4% (inconclusive) on its own test set, so the J1 contrast depends on the test

draw; J1 is not a refinement of the held path.

Cost. With the pre-specified protocol (inference mode, two warm-up and five timed calls, one batch of 256 sequences, 2 threads; Appendix B), native_S8_end_to_end/resampled_S8_end_to_end is 9.7–11.5 across seeds, forward-only 11.5–14.1, and on C0 0.98–1.04 (Table 7); end-to-end includes resampling (0.007–0.019 s on S8). The two boundaries are timed in separate jobs and cross in 4 of 60 cells (by up to 9.5%), so differences of a few percent are run-to-run variation.

What this does not show. Four seeds of one two-block model on one synthetic task with a teacher from the model’s own family replicate within that task only (Section 7); the primary estimand was chosen after the development run, and the calibration curves were unstable in every seed (up to 0.84 in the development run), so the selected checkpoints depend on the evaluation schedule.

5.2. Real sensor windows: UCI HAR, one checkpoint

Data and representation. We use the Human Activity Recognition Using Smartphones data (Anguita et al., 2013): 30 subjects, six activities, 3-axis total acceleration, body acceleration and gyroscope at 50 Hz, distributed as noise-filtered windows of 128 readings (2.56 s) with 50% overlap between consecutive windows of a subject; we use the nine inertial channels, not the 561 engineered features. These distributor-preprocessed windows are not a raw stream, and

the refinement below is artificial. The official subject partition is kept (21 training subjects, 9 test subjects, 2947 test windows); four training subjects, chosen before the run by a hash of their identifiers, are the calibration subjects (1364 windows), and the other 17 (5988 windows) fit the model and the per-channel normalization, which is never re-fitted. A window’s 128 values are represented as 128 held intervals of $1/50$ s ($T=2.56$ s), a chosen reconstruction: C0 keeps them; S8 splits every interval into 8 (1024 tokens); H8 splits the first 64 (576 tokens); held values, horizon and label are unchanged. Because windows overlap, every quantity is averaged within a subject first and the 9 subject values are averaged with equal weight.

Model, training and arms. The model is the official `TIDESCClassifier` of the pinned code with the backbone setting of Section 5.1 (23,786 parameters), nine input channels and the official head, a token mean of the features followed by a linear map to six classes. One model seed; Adam, batch 32, 2000 updates; the state with the lowest calibration cross-entropy (every 100 updates) is frozen and hashed before any test window is read: update 1200, calibration cross-entropy 0.092, accuracy 0.974. The same checkpoint is evaluated with three arms: the official token-mean readout on the condition’s grid (*native mean*); the same features pooled with weights dt_j/T at the same location and passed through the same affine head (*native time*, an inference-time ablation, not a trained baseline); and the official readout on the training grid rebuilt from the observed values (*resampled*), which on S8 and H8 is the identity control (input difference 0). Because the head $A(h) = Wh + b$ is affine and the weights sum to one, $A(\sum_j w_j h_j) = \sum_j w_j A(h_j)$ for both poolings (checked, Appendix D); softmax outputs are never averaged. The primary estimand, fixed before the run, is d_s , the mean over the windows of test subject s of $\text{CE}(\text{native mean}, \text{H8}) - \text{CE}(\text{native time}, \text{H8})$, and D , its equal-weight mean over the 9 subjects, with a subject-level paired percentile bootstrap (2000 draws); cross-entropy, accuracy, prediction flips, output distances and cost are secondary and reported separately.

Results (Tables 4 and 5). *Pooled splitting change and its common-time part.* Splitting every held interval (S8) changes the class-centered pooled logits by 0.029 (relative ℓ_2 ; subject range in Table 4), the class probabilities by a total variation of 0.008, and the predicted class of 0.008 of the windows, while the resampled arm reproduces C0 exactly; the cross-entropy change is -0.010 nats with a subject-level interval of $[-0.034, +0.012]$, so no risk change is established at nine subjects. Unlike the synthetic discrepancy of Section 5.1, which compares outputs at the same 32 times, this pooled change compares a mean over 128 time points with a mean over 1024, so it is not by itself a common-time (model-coupling) effect:

even an exact single layer has a split-invariant end value but a different token mean. A fixed-checkpoint diagnostic designed after these results (Table 12, Appendix D) separates the two with the official affine head A : $z_0 = A(\overline{H_0})$, $z_{\text{end}} = A(\text{mean}_k H_8[\text{end}(k)])$ from the S8 features at the 128 common time points, $z_{\text{all}} = A(\overline{H_8})$, and exactly $z_{\text{all}} - z_0 = (z_{\text{end}} - z_0) + (z_{\text{all}} - z_{\text{end}})$. The common-time change has relative norm 0.0018 and the extra-sampling change 0.0284 (per-window ratios $\|\Pi d\|_2 / \|\Pi z_0\|_2$ averaged subject-first, as in Table 5; as a ratio of subject-mean norms the common-time part is 0.0031), with inner product $+0.091 \logit^2$, so the two norms are not additive shares; the final features at the common end points move by 0.0025 relative, and z_{end} flips no prediction (cross-entropy 0.271 against 0.270 for z_0). The model’s outputs at the common times therefore do change under splitting, but by an order of magnitude less than the pooled readout; the extra-sampling term involves the model path through the added tokens and is not a pure quadrature error, and neither term is attributed to a layer. *Pooled weighting.* Under H8 the token-mean readout moves the centered logits by 0.210, changes the predicted class of 0.028 of the windows and the cross-entropy by $+0.012 [-0.005, +0.028]$; time weighting of the same features reduces these to 0.016, 0.003 and $-0.004 [-0.017, +0.007]$. The logit decomposition of Section 5.1 applies unchanged (norms in logit units, per window, then subject means): the token-mean change has norm 7.689, the density-weighting term from the C0 token logits alone 7.439, the model term M 0.997, its part on the 128 interval-end tokens 0.013, and the time-weighted change Q 0.565: as on the synthetic task, the large token-mean effect is mostly the weighted-average identity, and the model outputs at the base times change two orders of magnitude less. *Task risk.* The primary estimand is $D = +0.016$ nats with a subject-level interval of $[-0.004, +0.032]$; d_s is positive for 8 of the 9 subjects (range -0.044 to $+0.063$). The interval includes zero: the pre-registered difference is not established on this dataset, and non-significance is not evidence of equivalence; the sign pattern matches the synthetic finding, but nine clusters give a coarse interval. Accuracy is 0.904 on C0, 0.905 under H8 with the token mean and 0.903 with time weighting (no interval is reported). *Control and cost.* A control (time-weighted per-channel means and second moments with a multinomial linear classifier; not a task ceiling or a capacity-matched baseline) has identical features on C0, S8 and H8 (maximum difference 1.8×10^{-14} ; cross-entropy 0.695, accuracy 0.765), as any readout that integrates the held path over time must. Native evaluation costs `native_S8_end_to_end/resampled_S8_end_to_end` = 12.7 (forward-only 12.4) and 5.0 on H8 (Table 15).

Initialization sensitivity (two further model seeds). The frozen protocol was retrained with model seeds 101 and 102 and each checkpoint evaluated once on the *same* nine

Table 4. P1-HAR-01: one official TIDES classifier (one model seed) on UCI HAR inertial windows, per official test subject (288–381 windows each; windows overlap by 50%, so the subject is the analysis unit). C0: the 128 distributed values as 128 held intervals of 1/50 s; H8: the first 64 intervals split by 8 (576 tokens), same held values and labels. CE: mean cross-entropy (nats) of the window label under token-mean pooling (official head) and time-weighted pooling of the same features (inference ablation). d_s : primary difference CE(mean, H8) – CE(time, H8) per subject; the last row is the equal-weight mean over the 9 subjects with a subject-level paired percentile bootstrap interval (2000 draws), which is coarse with 9 clusters. acc.: accuracy (time-weighted pooling: 0.903 on H8). flip: fraction of windows whose predicted class under H8 (token mean) differs from C0.

subject	CE C0	CE H8 mean	CE H8 time	d_s	acc. C0	acc. H8 mean	flip H8 mean
2	0.193	0.217	0.194	+0.023	0.891	0.884	0.010
4	0.244	0.220	0.264	−0.044	0.896	0.909	0.050
9	0.455	0.469	0.432	+0.038	0.774	0.771	0.094
10	1.132	1.103	1.083	+0.020	0.714	0.714	0.054
12	0.106	0.110	0.100	+0.011	0.991	0.994	0.009
13	0.119	0.144	0.132	+0.012	0.954	0.954	0.012
18	0.072	0.087	0.075	+0.013	0.962	0.962	0.011
20	0.075	0.140	0.078	+0.063	0.952	0.955	0.008
24	0.037	0.049	0.038	+0.011	1.000	1.000	0.000
mean	0.270	0.282	0.266	+0.016 [−0.004, +0.032]	0.904	0.905	0.028

Table 5. P1-HAR-01 effects of the same-held-path refinement, equal-weight means over the 9 test subjects (subject means first). Δ CE: cross-entropy of the arm on the condition minus cross-entropy of the official token-mean readout on C0, with the subject-level bootstrap interval. flip: fraction of windows whose predicted class differs from C0. TV: total variation $\frac{1}{2} \sum_c |p_c - p_c^{C0}|$ between the class probabilities. c.-logit: $\|\tilde{z} - \tilde{z}^{C0}\|_2 / \|\tilde{z}^{C0}\|_2$ with class-centered logits (0 windows excluded by the 10^{-6} rule). The resampled arm reproduces the C0 input exactly (identity control). On S8 the grid is uniform, so the two native arms coincide.

cond.	arm	Δ CE vs. C0 [95% CI]	flip	TV	c.-logit
S8	native, token mean	−0.0097 [−0.034, +0.012]	0.008	0.008	0.029
S8	native, time-weighted	−0.0097 [−0.034, +0.012]	0.008	0.008	0.029
S8	resampled (identity control)	+0.0000 [+0.000, +0.000]	0.000	0.000	0.000
H8	native, token mean	+0.0119 [−0.005, +0.028]	0.028	0.037	0.210
H8	native, time-weighted	−0.0044 [−0.017, +0.007]	0.003	0.005	0.016
H8	resampled (identity control)	+0.0000 [+0.000, +0.000]	0.000	0.000	0.000

subjects (Appendix D): a sensitivity check on one architecture and one dataset, not an independent replication; the three runs are separate rows, never 27 observations, and the pre-registered primary of the original run stands. The separation repeats: the common-time part of the pooled S8 change is 0.0027 and 0.0034 against pooled changes of 0.0231 and 0.0445, z_{end} flips no prediction, and W dominates the H8 token-mean change in each run (Table 14). The risk numbers move more: $D = +0.050 [+0.023, +0.082]$ (8 of 9 subjects positive) and $+0.044 [+0.011, +0.077]$ (7 of 9) against +0.016, with H8 token-mean cross-entropy changes against C0 of +0.052 and +0.028 (time-weighted +0.002, −0.016): the direction is consistent across the three initializations, the size varies about three-fold, and the two new intervals exclude zero but share subjects and data with the original run and do not change its verdict.

What this does not show. One dataset and one architecture give results conditional on these checkpoints, not on the population of training runs, other sensors, naturally irregular acquisitions or TIDES’ benchmarks; nine overlapping-window subjects bound what the interval can resolve, the

common-time diagnostic is post hoc on the original checkpoint, and the two further seeds reuse its subjects and data.

6. Controlled Toy Results

These fixed-parameter results are numerical checks and controlled measurements, not proofs, and say nothing directly about trained models. The pre-registered FIRST_RUN used 32 independent units (one parameter and input draw each), $T=8$, $K=16$, $m \in \{1, 2, 4, 8\}$, 2 CPU threads and 20.6 s; details and deviations are in Appendix A.

Splitting. With $\Delta=dt\,g$ and exact ZOH the output is split-invariant to 6.3×10^{-15} and agrees with an independent RK4 reference to 1.6×10^{-13} (an engineering check of Proposition 4.1). Euler- B with the correct step converges at first order (median slope −1.03) but has a median base-step artifact of 0.21 in this toy; the magnitude depends on the toy’s $\Delta|\lambda|$ scale. Bilinear converges at second order (−2.02); with $\Delta=\tau g$ the artifact grows from 0.48 to 1.3 (medians, $m=2 \rightarrow 8$), by construction (Table 18, Figure 1a).




Figure 1. Error against the exact held-path solution as the same held interval is split into m sub-steps: (a) single layer (FR-E1-SPLIT; medians over 32 toy units); (b) two-layer cascade (P1-COMP-01; one deterministic case per line). Exact ZOH with $\Delta=dt\,g$ (a) and the exact coupled solution (b) stay at 10^{-15} and are not drawn. The $\Delta=\tau g$ line starts at $m=2$ because it is exact at $m=1$ by construction. Values are listed in Tables 18 and 19.

New observations. With exact ZOH the artifact against RK4 on the new held path stays below 1.7×10^{-13} , so the whole change is a response to the changed input reconstruction. The exact decomposition (2) of the other variants (exploratory, Appendix I) has cross terms between -1.37 and $+0.42$ of the per-unit total, so no per-unit share is defined.

Readouts and interaction of fixes. On a split-invariant backbone, a density change moves the sample mean by 0.30 and the exact integral by 1.5×10^{-16} (scale-normalized, $m=8$). The pre-registered clause H7c was mis-specified before the run and failed (slope $+0.42$); its restatement passed only weakly (median slope -0.076 , 21 of 32 units negative). One fix at a time (Table 17): exact ZOH alone gives 0.48, the physical step alone 0.30, a time readout alone 0.25, against 0.46 without fixes; in this tested ablation the three corrections combined reduced the artifact to 1.5×10^{-16} . This is not a necessity result: it concerns a time-average readout of one toy, a last-state readout needs no time integral, and special inputs or coefficients (for example a zero input) remove it without all three corrections. Under the pre-registered stop condition (an operational rule, not a statistical, equivalence or novelty test) the toy gives no support for a new architecture.

Two-layer cascade. For (3) (P1-COMP-01; held input of FIRST.RUN unit 0; 0.11 s), the exact coupled solution is split-invariant to 1.7×10^{-15} , while the layerwise computation has error 0.19 at the base step and 0.026 at $m=8$ (slope -0.95 ; -0.93 for $a=c$; Figure 1b, Table 19), consistent with, not a proof of, first order for this input.

7. Limitations

Synthetic evidence: four training seeds and a development seed of one small two-block model (one selective SSM) on one synthetic task whose teacher is from the model’s own family; it supports no statement about other tasks, the population of training runs, TIDES’ benchmarks, Mamba, S4 or S5, and the accuracy contrast was not stable. *Real-sensor evidence:* one dataset, one architecture, three model seeds on the same nine subjects (one pre-registered run and a two-seed sensitivity check), an artificial refinement of distributor-preprocessed windows (Section 5.2); the task-risk effect keeps its direction but its size varies across seeds, and no other model or dataset was evaluated. *Entry points not measured:* token-indexed convolution, normalization in training mode, and deeper or wider stacks. *Analysis:* Propositions 4.1 and 4.2 are standard; Proposition 4.2 is an upper bound for a linear time-invariant scalar cascade (Remark 4.3), with no lower bound, no statement about deep selective networks and no global hold-error or Euler-B bound. *Post hoc choices:* the primary estimand of the repetition was chosen after the development run; the layer localization, the H8 signed decomposition, the HAR common-time diagnostic, the toy decomposition and the stiffness columns were analysed after the results were seen.

8. Conclusion

Grid dependence in selective SSMs can enter through a step that counts samples, an inexact input-matrix discretization, sample-count readouts and, in stacks, the resampling of internal signals; the last is not caused by any change of the observed input. In the single-layer toy ablation only the combined corrections reduced the readout artifact to rounding level, and the inter-layer resampling error remained ($O(\text{mesh})$ upper bound for a scalar linear cascade). On the official TIDES code (physical step), splitting the held path changes the outputs at common times in all four training seeds of one synthetic task; on real sensor windows (UCI HAR) the common-time part of the pooled splitting change is an order of magnitude smaller than the readout term over the added tokens (not additive shares) in all three model seeds; the larger readout effect of a same-path density change is, in both, mostly the weighted-average identity of token-mean pooling, which time weighting reduces to a small residual; and the loss consequence does not replicate across synthetic seeds, while on the sensor windows the pre-registered subject-level interval includes zero and two further seeds give larger positive values on the same subjects. Whether these effects matter for other selective SSMs, for naturally irregular acquisitions, and across training runs is open.

Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning, specifically the understanding of how sequence models respond to changes in sampling. The experiments use synthetic data and one public, anonymized activity-recognition dataset (UCI HAR, distributed for research with a citation requirement), all on small CPU computations; the trained-model results come from four training seeds and a development seed on one synthetic task and from one checkpoint on the public dataset. No clinical or otherwise sensitive data were used. The main risk we see is over-reading such results: a model that behaves consistently in a controlled setting may still behave differently on irregularly sampled real data, for example clinical or sensor streams, so any use in such settings needs evaluation on the actual data and sampling process. We are not aware of other societal consequences that must be specifically highlighted here.

References

- Anguita, D., Ghio, A., Oneto, L., Parra, X., and Reyes-Ortiz, J. L. A Public Domain Dataset for Human Activity Recognition Using Smartphones. In *21st European Symposium on Artificial Neural Networks, Computational Intelligence and Machine Learning (ESANN)*, 2013. Dataset: UCI Machine Learning Repository, DOI 10.24432/C54S4K.
- Cirone, N. M., Orvieto, A., Walker, B., Salvi, C., and Lyons, T. Theoretical Foundations of Deep Selective State-Space Models. arXiv:2402.19047, 2024.
- Dao, T. and Gu, A. Transformers are SSMs: Generalized Models and Efficient Algorithms Through Structured State Space Duality. arXiv:2405.21060, 2024.
- Gu, A. and Dao, T. Mamba: Linear-Time Sequence Modeling with Selective State Spaces. arXiv:2312.00752, 2023.
- Gu, A., Goel, K., and Ré, C. Efficiently Modeling Long Sequences with Structured State Spaces. arXiv:2111.00396, 2021a.
- Gu, A., Johnson, I., Goel, K., Saab, K., Dao, T., Rudra, A., and Ré, C. Combining Recurrent, Convolutional, and Continuous-time Models with Linear State-Space Layers. arXiv:2110.13985, 2021b.
- Gu, A., Gupta, A., Goel, K., and Ré, C. On the Parameterization and Initialization of Diagonal State Space Models. arXiv:2206.11893, 2022.
- Hasani, R., Lechner, M., Wang, T.-H., Chahine, M., Amini, A., and Rus, D. Liquid Structural State-Space Models. arXiv:2209.12951, 2022.
- Kidger, P., Morrill, J., Foster, J., and Lyons, T. Neural Controlled Differential Equations for Irregular Time Series. arXiv:2005.08926, 2020.
- Lahoti, A., Li, K. Y., Chen, B., Wang, C., Bick, A., Kolter, J. Z., Dao, T., and Gu, A. Mamba-3: Improved Sequence Modeling using State Space Principles. arXiv:2603.15569, 2026.
- Nguyen, E., Goel, K., Gu, A., Downs, G. W., Shah, P., Dao, T., Baccus, S. A., and Ré, C. S4ND: Modeling Images and Videos as Multidimensional Signals Using State Spaces. arXiv:2210.06583, 2022.
- Rusch, T. K. and Rus, D. Oscillatory State-Space Models. arXiv:2410.03943, 2024.
- Smith, J. T. H., Warrington, A., and Linderman, S. W. Simplified State Space Layers for Sequence Modeling. arXiv:2208.04933, 2022.
- Soydan, T., Bessa, M. A., Mohr, D., and Barreira, R. TIDES: Implicit Time-Awareness in Selective State Space Models. arXiv:2605.09742, 2026.
- Walker, B., McLeod, A. D., Qin, T., Cheng, Y., Li, H., and Lyons, T. Log Neural Controlled Differential Equations: The Lie Brackets Make a Difference. arXiv:2402.18512, 2024. arXiv journal-ref: Proceedings of the 41st International Conference on Machine Learning, 2024.
- Zubić, N., Gehrig, M., and Scaramuzza, D. State Space Models for Event Cameras. arXiv:2402.15584, 2024. arXiv journal-ref: IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Seattle, 2024.

A. Pre-registration, deviations and compute

FIRST_RUN. Hypotheses H1–H8, metrics, 32 units, 2 CPU threads, a 120 s budget and stop conditions were committed before any experiment code ran. Amendments recorded before any experiment ran: A1 (adaptive RK4 reference step), A2 (a restated H7c, keeping the mis-specified original), A3 (secondary configurations fixed to seeds 0–7, descriptive). Deviations: a timing-only smoke test on seed 999 (metrics not inspected); a repeated execution after a manifest bookkeeping bug, with byte-identical data files; a partial execution killed by a shell pipe, whose outputs were deleted. P1-COMP-01 was pre-registered separately before its code ran. Table 6 lists every pre-registered verdict, including the failure. *Literature:* the key equations and experiment sections of TIDES, Mamba-3 and S5 were read in the original text; other entries were checked by keyword search.

P1-REAL-01. The design, the pinned implementation and the predictions were committed before any installation. Amendments R-A1–R-A9 were committed after an untrained adapter check and a timing smoke on four training trajectories, and before training: calibration-based checkpoint selection, same-held-path refinements only (new-observation conditions removed, one random-time condition kept as secondary), a resampler restricted to observed values, the pooled readout and pooling-only arm, the explicit equivalence unit, the single-seed pilot label, time caps, and seeding of all random generators (the TIDES initialization also draws from NumPy’s global generator). The pre-registered predictions were: native splitting discrepancy at S8 above 10^{-4} (supported), resampled splitting discrepancy exactly zero (supported), native runtime growing with tokens (observed), and a larger sample-mean than time-weighted pooled change on H8 (supported).

P1-REAL-02. The configuration (seeds, identical protocol, new test set, primary and secondary estimands, bootstrap and margin rules, timing boundaries, failure handling, conditional extra diagnostics) was committed before any repetition run. One incident occurred: the evaluation of the development row crashed before writing any output because its output directory was not created; the attempt is recorded as INVALID, the directory creation was added, and only that row was rerun (the four training seeds were unaffected). The pre-specified conditional layer diagnostic was not triggered.

H8 signed decomposition (derived, after the results). P1-REAL-02 stored per-trajectory pooled values but not per-token outputs. A configuration fixed before execution specified one frozen native forward pass per existing checkpoint on C0 and H8 (no training, no timing, same test set), a check that it reproduces the stored pooled values (maximum absolute difference 0), and the signed identities of Table 2 (largest residual 4.6×10^{-15}).

Compute. All on CPU (2 threads for experiments), no GPU, no paid API; the synthetic stages use no external data, and the only external data are the public UCI HAR windows of Appendix D; wall and process CPU seconds are distinguished where recorded. FIRST_RUN: two executions of about 21 s, one partial execution and a smoke test of about 3 s together. P1-COMP-01: 0.11 s. P1-REAL-01 (development): installation of a CPU tensor library (35 s) and of a $\LaTeX$ toolchain (66 s); adapter check and timing smoke (≈ 4 s; 0.031 s per batch-32 step); convention check (≈ 2 s); training including label generation (97 s); evaluation (24 s); layer diagnostic (≈ 11 s). P1-REAL-02: data generation 40 s; seed-0 equivalence retrain ≈ 57 s; four training runs 188 s wall in total (364 s process CPU in the training loops); test-set generation 16 s; five evaluations 143 s wall in total, plus the invalid development attempt (≈ 30 s); H8 decomposition forward passes 7.7 s.

Table 6. Pre-registered check verdicts, including the failed H7c_original (mis-specified before the run; amendment A2) and the weakly supported H7c_restated. These are bookkeeping, not scientific results.

run	check	verdict
FIRST_RUN	H1	PASS
FIRST_RUN	H2	PASS
FIRST_RUN	H3	PASS
FIRST_RUN	H4	PASS
FIRST_RUN	H5a	PASS
FIRST_RUN	H5b	PASS
FIRST_RUN	H6	PASS
FIRST_RUN	H7a	PASS
FIRST_RUN	H7b	PASS
FIRST_RUN	H7c_original	FAIL
FIRST_RUN	H7c_restated	PASS (weak; see text)
FIRST_RUN	H8	PASS
P1-COMP-01	COMP-H1 (distinct)	PASS
P1-COMP-01	COMP-H2 (distinct)	PASS
P1-COMP-01	COMP-H3 (distinct)	PASS
P1-COMP-01	COMP-H4 (distinct)	PASS
P1-COMP-01	COMP-H5 (distinct)	PASS
P1-COMP-01	COMP-H1 (equal)	PASS
P1-COMP-01	COMP-H2 (equal)	PASS
P1-COMP-01	COMP-H3 (equal)	PASS
P1-COMP-01	COMP-H4 (equal)	PASS
P1-COMP-01	COMP-H5 (equal)	PASS

B. Implementation and protocol details

Pinned code facts (read, then executed through the adapter): the step is `step_scale * exp(log_step)` with a per-mode timescale; the ZOH input matrix is $(1/\text{Lambda}) * (\text{Lambda_bar}-1) * B$ without `expm1`; the scan is $x_k = \bar{\Lambda}_k x_{k-1} + \bar{B}_k u_k$ with the gap before step k ; each block applies `BatchNorm1d(affine=False)` over batch and time (running statistics at evaluation); an optional causal convolution defaults to off; the classifier head pools with `h.mean(dim=1)`. *Adapter check (untrained, four training trajectories):* forward outputs finite; gradients finite for all 54 parameters that receive one (the six static Re Λ , B , C tensors are unused in input-dependent mode; parameters upstream of the zero-initialized head projections receive zero gradient at step 0 only); an Adam step changed the parameters and lowered the loss on the same batch; evaluation was deterministic, and a per-step step tensor of ones matched the scalar step exactly. *Labels:* teacher output on the continuous path from an independent RK4 integration (128 sub-steps per training interval); the pooled label is its time average by Simpson’s rule. *Readout of J1:* the last native output at or before each training-grid time. *Eval-mode checks (P1-REAL-02, every seed):* every module in evaluation mode during calibration and test; dropout probability 0; batch normalization with running statistics; equal sequence lengths within every batch (no padding or masking is used); all step sizes positive (the smallest native step is 1/8 of the training step on S8 and H8). *Timing (P1-REAL-02):* Table 7.

Table 7. P1-REAL-02 timing (one batch of 256 sequences; median of 5 timed calls after 2 warm-ups; `inference_mode`; 2 threads). End-to-end includes tensor construction, LOCF resampling for the resampled arm, forward, de-normalization and readout mapping; forward-only is the model call on prebuilt tensors. CPU seconds are process CPU time summed over threads. Timed repeats are not independent samples. [‡]Development checkpoint (seed 0).

seed	cond.	tokens	nat. e2e	nat. fwd	res. e2e	res. fwd	nat. e2e CPU
0 [‡]	C0	32	0.049	0.044	0.047	0.042	0.10
0 [‡]	S8	256	0.578	0.574	0.068	0.063	1.12
0 [‡]	H8	144	0.297	0.281	0.069	0.056	0.58
1	C0	32	0.042	0.039	0.042	0.041	0.09
1	S8	256	0.565	0.545	0.049	0.041	1.10
1	H8	144	0.308	0.275	0.047	0.042	0.60
2	C0	32	0.041	0.038	0.042	0.038	0.08
2	S8	256	0.560	0.550	0.058	0.039	1.09
2	H8	144	0.300	0.281	0.053	0.043	0.59
3	C0	32	0.051	0.044	0.049	0.043	0.10
3	S8	256	0.544	0.566	0.054	0.047	1.06
3	H8	144	0.286	0.275	0.061	0.047	0.56
4	C0	32	0.045	0.041	0.045	0.042	0.09
4	S8	256	0.574	0.531	0.054	0.046	1.10
4	H8	144	0.281	0.276	0.055	0.047	0.56

C. Development pilot (seed 0) on its own test set

Table 8 gives the development checkpoint on the new shared test set. The remaining results come from the development run (P1-REAL-01), whose questions and estimands were not fixed in advance beyond its own predictions, on its own test set (identifiers 640–895). Its equivalence analysis used a margin of 5% of the point-estimated resampled MSE (kept here as exploratory); the repetition in Section 5 instead bootstraps the ratio r_s . Pooled quantities in Table 10 are mean absolute errors (or changes) of pooled predictions against the pooled label, in label units; they are not hidden-state discrepancies or MSEs. Its runtimes are forward-only model calls on prebuilt tensors (no warm-up, outside inference mode) and exclude resampling. *Layer diagnostic (exploratory, not pre-registered)*. After seeing the development results we measured the change at intermediate features of the development checkpoint (Table 11): the encoder output is unchanged, the first block changes by at most 7.2×10^{-6} in float32 and 1.8×10^{-14} with the weights cast to float64, and the second block by 4.1×10^{-3} to 7.2×10^{-3} in both precisions. The first block receives the held input and uses an exact held-coefficient update, whereas the second block receives the first block’s output, which splitting samples at more points within each interval; neither proposition is a statement about this trained network. The repetition pre-specified a per-seed layer diagnostic only if some seed’s S8 discrepancy fell below 10^{-3} ; this did not happen, so the localization rests on one checkpoint.

Table 8. Development checkpoint (seed 0, P1-REAL-01) evaluated with the P1-REAL-02 protocol on the new shared test set; not part of the replication criterion. Definitions as in Tables 1 to 3.

quantity	development checkpoint, new test set
selected step / calibration MSE	1500 / 0.0119
S8 disc. [95% CI]	0.0114 [0.0108, 0.0120]
D_s [95% CI]	0.575 [0.511, 0.643]
$p^{\text{mean}} - p(\text{C0}), W, M, Q$ (means)	+0.452, +0.460, −0.009, −0.009
$r_s(\text{S2})$ [95% CI], verdict	−2.9 [−4.6, −1.5] =
$r_s(\text{S8})$ [95% CI], verdict	−3.9 [−6.7, −1.8] <
$r_s(\text{H8})$ [95% CI], verdict	−6.5 [−9.4, −3.8] <
$r_s(\text{J1})$ [95% CI], verdict	−20.9 [−30.8, −8.0] ≪

Table 9. DEVELOPMENT run (P1-REAL-01): trained official TIDES, seed 0, on its own test set (256 trajectories, ids 640–895; one checkpoint selected on the calibration split; paired). “disc.”: mean relative ℓ_2 change of the per-step outputs at the 32 training-grid times against the same arm on C0. MSE: per-trajectory mean squared error to the teacher at those times (original units; label s.d. 2.35). Δ MSE: mean paired difference native – resampled in units of 10^{-4} (label units squared) with a 95% bootstrap CI; the margin is $\pm 5\%$ of the point-estimated resampled mean MSE of the same condition (exploratory rule of the development run). Verdict (native vs. resampled MSE): lower/higher if the CI excludes 0 on that side. Runtimes of this run are not tabulated; they were forward-only and are superseded by Table 7. C0 is the training grid; S m and H8 (first half split by 8) keep the held path; † J1 (random observation times) is a secondary non-refinement condition whose held path differs by construction.

cond.	disc. native	disc. resampled	MSE native	MSE resampled	Δ MSE [95% CI] (10^{-4})	margin rule
C0	0	0	0.0167	0.0167	0 [0, 0]	identical inputs (check)
S2	6.1×10^{-3}	0	0.0162	0.0167	-5.1 [-7.3, -2.7]	within $\pm 5\%$
S4	9.2×10^{-3}	0	0.0161	0.0167	-6.5 [-9.9, -2.9]	lower
S8	0.011	0	0.0160	0.0167	-6.9 [-10.8, -2.7]	lower
H8	7.8×10^{-3}	0	0.0156	0.0167	-10.8 [-14.0, -7.6]	lower
J1 †	0.35	0.39	1.14	1.22	-785 [-2120, +623]	inconclusive

Table 10. Pooled readout of the development checkpoint on its own test set (P1-REAL-01), label units, means over 256 test trajectories. Left: $n^{-1} \sum_i |p_i(c) - \bar{y}_i|$; right: $n^{-1} \sum_i |p_i(c) - p_i(C0)|$, with $p = p^{\text{mean}}$ (“native mean”, token mean as in the TIDES classifier head), $p = p^{\text{time}}$ (“native time”, right-endpoint time-weighted sum of the same native outputs; a pooling-only ablation, not an exact integral) or $p = p^{\text{res}}$ (“resampled”, token mean of the 32 resampled outputs); definitions in Section 5. Pooled label $\bar{y}_i$: time average of the teacher output over $[0, T]$ by Simpson’s rule on its RK4 nodes (mean absolute value 1.59).

cond.	error to pooled label			change vs. C0		
	native mean	native time	resampled	native mean	native time	resampled
C0	0.063	0.063	0.063	0	0	0
S8	0.070	0.070	0.063	0.032	0.032	0
H8	0.56	0.068	0.063	0.55	0.016	0
J1	0.37	0.28	0.29	0.36	0.26	0.30

Table 11. Where the splitting change appears in the development checkpoint (EXPLORATORY, added after its results were seen). Mean relative ℓ_2 change against C0 at the training-grid times of the features after each block and of the (normalized) output, with the frozen weights evaluated in float32 and cast to float64. The encoder output is unchanged (0) in all cases.

cond.	block 1		block 2		output	
	fp32	fp64	fp32	fp64	fp32	fp64
S2	5.6×10^{-6}	5.1×10^{-15}	4.1×10^{-3}	4.1×10^{-3}	6.8×10^{-3}	6.8×10^{-3}
S4	6.0×10^{-6}	1.1×10^{-14}	6.1×10^{-3}	6.1×10^{-3}	0.010	0.010
S8	7.2×10^{-6}	1.8×10^{-14}	7.2×10^{-3}	7.2×10^{-3}	0.012	0.012
H8	5.8×10^{-6}	1.4×10^{-14}	5.5×10^{-3}	5.4×10^{-3}	9.1×10^{-3}	9.1×10^{-3}

D. P1-HAR-01: data, protocol and cost

Data. The official archive (61,005,872 bytes; sha256 recorded with the per-file hashes in the acquisition record) was downloaded once from the UCI repository, checked for path traversal before extraction, and kept outside version control. The repository page states a CC BY 4.0 license; the README requires citing [Anguita et al. \(2013\)](#) and prohibits commercial use; both statements are recorded. Channels: total acceleration (g), body acceleration (g, gravity removed by the distributors' 0.3 Hz filter) and angular velocity (rad/s). Train 7352 windows / test 2947 windows, as distributed. *Calibration rule.* The 21 training subjects are ranked by the SHA-256 digest of "P1-HAR-01:calibration:<id>" and the first four (7, 15, 17, 19) are the calibration subjects; the runner re-derives the rule. *Budget rule.* A 20-update smoke run (model discarded; run twice, before and after the pre-execution amendment, 16 s and 13 s, both in the ledger) measured the per-update and per-token cost and projected the training and evaluation cost; the pre-registered rule allowed the planned 2000 updates. Ledger: training 183 s wall (354 s process CPU), evaluation 170 s wall; all stages together 388 s wall and 719 s process CPU on 2 threads, against a cap of 3600 s / 7200 s. *Checks.* Eval mode enforced at calibration and test (no dropout, running batch statistics); all step scales positive; no padding or mask; the resampled input equals C0 on S8 and H8 to 0; the head-of-mean identity $A(\sum_j w_j h_j) = \sum_j w_j A(h_j)$ holds on C0 to 2.3×10^{-5} and the two native arms coincide on the uniform C0 grid to 1.5×10^{-5} ; the identities of the logit decomposition hold to 2.7×10^{-5} ; 0 windows excluded by the centered-logit rule. *Timing (Table 15).* native_S8_end_to_end / resampled_S8_end_to_end = 12.7; forward-only 12.4; H8 end-to-end 5.0. *Common-time diagnostic (P1-HAR-CT-01; exploratory, fixed checkpoint, designed after the P1-HAR-01 results).* The stored per-token logits (float16) reproduce the stored pooled logits only to 7.0×10^{-3} , so one float32 streaming forward of C0 and S8 over the 2947 test windows was run (batches of 256, eval mode, no intermediate tensors stored); it reproduces the stored C0 and S8 pooled logits to 0. The 128 common end points are the S8 tokens $8k$ ($k = 1..128$), verified against the time stamps. The split is an algebraic identity of the three logit vectors (the recorded residual, 0, is zero by construction and checks arithmetic only, not the model); the class centering $\Pi = I - \mathbf{1}\mathbf{1}^\top / 6$ is applied to every term alike, and the relative norms are per-window ratios averaged subject-first (as a ratio of subject-mean norms the common-time part is $0.093/30.4 = 0.0031$). The pre-specified float64 check (18 hash-selected windows) was *NOT NEEDED* by its rule (threshold 10^{-3} on the subject-mean relative common-time change; observed 0.0018); in absolute terms the common-time change (0.093 logits, subject mean) is about 3,042 times the float32 pooling residual recorded by P1-HAR-01. The common-time term is the change of the pooled readout over the 128 end points, not a per-token change: a per-token common-time logit change was not computed (the per-end-point feature change, 0.0025 relative, at most 0.0054 in a subject, is recorded but does not establish the per-token logit change). Cost: 73 s wall, 115 s process CPU on 2 threads (timer started before importing the framework). Norm sums and norm ratios are not contributions; the inner product is reported alongside.

Initialization sensitivity (R8-P1-HAR-REPLICATION). The contract (seeds, stages, inputs with hashes, estimands, budget) was committed before execution. Seeds 101 and 102 were never used before in this project; seed s seeds the parameter initialization (python, NumPy and torch generators) and a separate minibatch-order generator; data, splits and normalization are fixed arrays. Each checkpoint was evaluated once per condition in batches of 256 windows; the S8 common-end-point outputs and the H8 per-token logits were taken from the same forward passes and stored in float64; the resampled arm's input equals C0 exactly (checked), so its logits are the C0 logits and no forward was spent; timing was not re-measured. Selected updates 1600 and 1600 (calibration cross-entropy 0.120, 0.083). Cost: 674 s wall and 1281 s process CPU for the five stages on 2 threads (caps 1800 s / 3600 s); the first training stage ran at 280 s wall under concurrent CPU load from bookkeeping commands, within its 300 s allowance. Tables 13 and 14 give the three runs as separate rows.

Table 12. P1-HAR-CT-01 (exploratory, fixed checkpoint, designed after the P1-HAR-01 results): decomposition of the pooled S8 change of the HAR checkpoint, per test subject (subject means first; last row: equal-weight mean). With the official affine head A , $z_0 = A(\overline{H_0})$ (C0 readout), $z_{\text{end}} = A(\text{mean}_k H_s[\text{end}(k)])$ (S8 features at the 128 common time points), $z_{\text{all}} = A(\overline{H_8})$ (S8 readout), and $z_{\text{all}} - z_0 = (z_{\text{end}} - z_0) + (z_{\text{all}} - z_{\text{end}})$ per window exactly. Columns 2–4: per-window ratios $\|\Pi d\|_2 / \|\Pi z_0\|_2$ of the class-centered terms (all / common-time / extra-sampling), averaged within subject and then over subjects, the same definition as the centered-logit distance of Table 5; 5: inner product of the two centered terms (logit^2 ; the norms are not shares and do not add); 6: mean relative change of the final features at the common end points; 7–9: cross-entropy of z_0 , z_{end} , z_{all} ; 10–11: fraction of windows whose predicted class differs from z_0 . No interval and no test: a descriptive diagnostic of one checkpoint.

subject	rel. all	rel. common	rel. extra	$\langle \cdot, \cdot \rangle$	feat. end	CE z_0	CE z_{end}	CE z_{all}	flip end	flip all
2	0.0332	0.0024	0.0321	+0.490	0.0035	0.193	0.193	0.196	0.000	0.003
4	0.0224	0.0015	0.0217	+0.029	0.0015	0.244	0.244	0.280	0.000	0.009
9	0.0313	0.0013	0.0307	+0.125	0.0019	0.455	0.454	0.412	0.000	0.017
10	0.0295	0.0019	0.0288	+0.171	0.0029	1.132	1.133	1.037	0.000	0.031
12	0.0331	0.0020	0.0324	+0.013	0.0027	0.106	0.106	0.092	0.000	0.000
13	0.0309	0.0017	0.0301	+0.100	0.0019	0.119	0.119	0.138	0.000	0.009
18	0.0251	0.0010	0.0247	+0.009	0.0010	0.072	0.072	0.079	0.000	0.003
20	0.0256	0.0036	0.0253	−0.156	0.0054	0.075	0.076	0.074	0.000	0.003
24	0.0303	0.0010	0.0299	+0.038	0.0014	0.037	0.037	0.038	0.000	0.000
mean	0.0290	0.0018	0.0284	+0.091	0.0025	0.270	0.271	0.261	0.000	0.008

Table 13. Initialization sensitivity (R8-P1-HAR-REPLICATION): the frozen P1-HAR-01 protocol trained with two further model seeds (101, 102) on the same training data and evaluated on the SAME nine test subjects; the original run (seed 0) is repeated from Tables 4 and 5. Rows are separate runs and are never pooled (three runs on the same subjects are not 27 independent observations). step / calib. CE: selected update and its calibration cross-entropy. CE C0, acc. C0: subject-first means of the official readout on C0. D : primary difference $\text{CE}(\text{mean}, H_8) - \text{CE}(\text{time}, H_8)$, equal-weight mean over subjects, with the per-run subject-level paired bootstrap interval (2000 draws; coarse with 9 clusters) and the number of subjects with $d_s > 0$. ΔCE : cross-entropy change against C0 for S8 (token mean) and H8 (token mean / time-weighted).

seed	step	calib. CE	CE C0	acc. C0	D [95% CI]	$d_s > 0$	$\Delta\text{CE S8}$	$\Delta\text{CE H8, mean}$	$\Delta\text{CE H8, time}$
0 (orig.)	1200	0.092	0.270	0.904	+0.016 [−0.004, +0.032]	8/9	−0.010	+0.012	−0.004
101	1600	0.120	0.252	0.916	+0.050 [+0.023, +0.082]	8/9	+0.005	+0.052	+0.002
102	1600	0.083	0.252	0.923	+0.044 [+0.011, +0.077]	7/9	−0.026	+0.028	−0.016

Table 14. Initialization sensitivity of the decompositions (same three runs as Table 13; subject-first means; no interval, no test). S8: class-centered relative change of the pooled logits (“all”), its common-time part $\|\Pi(z_{\text{end}} - z_0)\| / \|\Pi z_0\|$ and extra-sampling part (per-window ratios, subject-first; not additive shares), the inner product of the two centered terms (logit^2), and the fraction of windows whose predicted class differs from C0 under z_{all} and under z_{end} . H8: class-centered relative change of the token-mean and time-weighted readouts, and the norms (logit units) of the token-mean change, its density-weighting term W , the model term M , its interval-end part M_{end} and the time-weighted change Q .

seed	S8						H8						
	all	common	extra	$\langle \cdot, \cdot \rangle$	flip all	flip end	mean	time	$\ \Delta\ $	$\ W\ $	$\ M\ $	$\ M_{\text{end}}\ $	$\ Q\ $
0 (orig.)	0.0290	0.0018	0.0284	+0.091	0.008	0.000	0.210	0.016	7.69	7.44	0.997	0.013	0.565
101	0.0231	0.0027	0.0216	+0.215	0.003	0.000	0.242	0.012	9.99	9.60	0.712	0.031	0.457
102	0.0445	0.0034	0.0451	−0.142	0.006	0.000	0.204	0.023	5.92	5.77	1.366	0.016	0.771

Table 15. P1-HAR-01 timing (one batch of the first 256 test windows; median of 5 timed calls after 2 warm-ups; inference_mode: 2 threads). End-to-end includes tensor construction, LOCF resampling for the resampled arm, forward and pooling; forward-only is the model call on prebuilt tensors. CPU seconds are process CPU time summed over threads. Timed repeats are not independent samples.

cond.	tokens	nat. e2e	nat. fwd	res. e2e	res. fwd	nat. e2e CPU
C0	128	0.276	0.306	0.282	0.353	0.55
S8	1024	3.562	3.288	0.281	0.266	6.52
H8	576	1.359	1.491	0.273	0.273	2.55

E. Interval conventions in the TIDES paper and code: details

The TIDES paper writes $x_{k+1} = \bar{\Lambda}(\Delta_k)x_k + \bar{B}(\Delta_k)u_k$ with $\Delta_k = t_{k+1} - t_k$ and $y_k = Cx_k + Du_k$, i.e. u_k is held forward on $[t_k, t_{k+1})$ and the state in y_k does not yet contain u_k (Soydan et al., 2026). The pinned code scans $z_k = \bar{\Lambda}(t_k - t_{k-1})z_{k-1} + \bar{B}(t_k - t_{k-1})u_k$ with $y_k = \text{Re}(Cz_k) + Du_k$, and its forecasting collate builds the step sequence $[t_0, t_1 - t_0, \dots]$: u_k is held backward on $(t_{k-1}, t_k]$ and the state in y_k contains it. Both are exact ZOH solutions, but of different input reconstructions. On a uniform grid every gap is equal, so $z_k = x_{k+1}$: the difference is a one-index shift of the state, which changes the alignment between the state part of y_k and the feedthrough Du_k . On an irregular grid each value is paired with a different gap, so the difference is not a relabeling. With the pinned code, untrained weights and a fixed impulse input, the code equals the backward-hold recurrence to 6.5×10^{-19} ; the one-index shift reproduces the paper form to 8.7×10^{-19} on a uniform grid but leaves 3.7×10^{-3} on an irregular one. The released code, not the printed equations, defines the model behind TIDES’ reported results, and it coincides with the right-endpoint convention used throughout this paper; our splitting protocol repeats the right-endpoint value accordingly. We make no claim that this affects TIDES’ results.

F. Proof of Proposition 4.2

Proof. (a) follows as in Proposition 4.1, since the coupled system has constant coefficients on each held interval. (b) Layer 1 is exact, so with $\tau_k = dt_k$, $e_k = e^{-c\tau_k}e_{k-1} + d \int_0^{\tau_k} e^{-c(\tau_k-s)}[h_1(t_k) - h_1(t_{k-1} + s)] ds$. If $a > 0$ then $|h_1| \leq |b|U/a$, so $a|h_1| \leq |b|U$; if $a = 0$ then $a|h_1| = 0$; either way $|\dot{h}_1| \leq 2|b|U$. On each held interval h_1 is C^1 , so the bracket is at most $2|b|U(\tau_k - s)$, and with $e^{-c(\tau_k-s)} \leq 1$ the local term is at most $|b|d|U\tau_k^2$. Unrolling from $e_0 = 0$ gives $|e_k| \leq |b|d|U\delta \sum_{j \leq k} \tau_j e^{-c(t_k-t_j)}$, and since $e^{-c(t_k-t_j)} \leq e^{c\delta} e^{-c(t_k-s)}$ for $s \in (t_{j-1}, t_j]$, the sum is at most $e^{c\delta} \int_0^{t_k} e^{-c(t_k-s)} ds \leq e^{c\delta}/c$. $\square$

In the stored P1-COMP-01 outputs, the largest error at the base times is 9–20% of the bound for $m \in \{1, 2, 4, 8\}$, which is consistent with, not a proof of, the bound.

G. Numerical floor

The naive coefficient $(e^z - 1)/\lambda$ loses relative precision as $|z|$ shrinks (8.0×10^{-4} at $z = -10^{-14}$ in float64), while `expm1` stays at rounding level (Table 16). In emulated float32, split composition of one scalar mode drifts to 1.1×10^{-4} at 16,384 sub-steps, so differences below about 10^{-4} are not interpretable in that emulation; this is not a floor for a trained multi-layer network with 256 tokens, for which Section 5 compares float32 with float64 directly.

Table 16. Numerical floor (FR-E4). Left: relative error of the ZOH input coefficient $(e^z - 1)/\lambda$ against a 50-digit reference. Right: splitting $[0, 1]$ into m held sub-steps ($\lambda = -1$, $g = \ln 2$, exact $h(1) = 0.5$); float32 is an *emulation* (every primitive rounded, correctly rounded exp), not a GPU kernel.

z	naive	expm1	m	fp64 expm1	fp32 expm1	fp32 naive
-0.01	5.4×10^{-15}	4.7×10^{-17}	1	0	0	0
-1×10^{-4}	1.3×10^{-14}	3.4×10^{-17}	4	2.2×10^{-16}	0	6.0×10^{-8}
-1×10^{-6}	1.6×10^{-11}	7.9×10^{-17}	16	2.2×10^{-16}	1.2×10^{-7}	2.4×10^{-7}
-1×10^{-8}	1.1×10^{-9}	6.4×10^{-17}	64	4.4×10^{-16}	6.0×10^{-8}	0
-1×10^{-10}	8.3×10^{-8}	3.4×10^{-17}	256	1.8×10^{-15}	4.3×10^{-6}	5.0×10^{-6}
-1×10^{-12}	2.2×10^{-5}	2.4×10^{-17}	1024	4.4×10^{-15}	1.8×10^{-5}	2.0×10^{-5}
-1×10^{-14}	8.0×10^{-4}	4.9×10^{-17}	4096	1.1×10^{-13}	2.1×10^{-5}	2.8×10^{-5}
			16384	5.2×10^{-13}	1.1×10^{-4}	2.3×10^{-4}

H. Reproducibility

The anonymized supplementary package contains the code (toy, RK4 references, cascade check, TIDES adapter, convention check, runners, diagnostics and the generator that produces every table, figure and number in this paper), the configurations, and the raw records from which the generator works, together with the sha256 of every checkpoint. It contains the 116 KB UCI HAR checkpoint with its SHA-256, but not the synthetic-task checkpoints (their SHA-256 values are listed and they are regenerated by the training command in the package; a retrain of the development seed reproduced every tensor), not the

Table 17. One fix at a time (FR-E3-READOUT, split of the first half, $m=8$; median scale-normalized readout change over 32 units). The effects are not additive, so no per-fix share is reported.

fixes applied	change
none ($\Delta=\tau g$, Euler- B , sample mean)	0.46
time readout only (trapezoidal)	0.25
exact ZOH B only	0.48
$\Delta=dt\ g$ only	0.30
$\Delta=dt\ g$ + trapezoidal	0.049
$\Delta=dt\ g$ + exact ZOH (sample mean)	0.30
$\Delta=dt\ g$ + exact ZOH + trapezoidal	0.011
$\Delta=dt\ g$ + exact ZOH + exact integral	1.5×10^{-16}

Table 18. Splitting a held interval into m sub-steps (FR-E1-SPLIT; toy, 32 units, uniform base grid, real modes). Entries are the median over units of the output artifact $\|y_m - y^*\|_2 / \|y^*\|_2$ at the 17 common base times, where y^* is the exact held-path solution; slope is the median per-unit log-log slope over $m \in \{1, 2, 4, 8\}$. τ is the base step, so $\Delta=\tau g$ coincides with $\Delta=dt\ g$ at $m=1$.

variant	$m=1$	$m=2$	$m=4$	$m=8$	max, $m=8$	slope
exact ZOH, $\Delta=dt\ g$	0	2.6×10^{-16}	3.8×10^{-16}	5.5×10^{-16}	6.3×10^{-15}	–
Euler- B , $\Delta=dt\ g$	0.21	0.10	0.050	0.025	0.12	–1.03
bilinear, $\Delta=dt\ g$	9.5×10^{-3}	2.3×10^{-3}	5.6×10^{-4}	1.4×10^{-4}	8.1×10^{-4}	–2.02
exact ZOH, $\Delta=\tau g$	0	0.48	0.93	1.3	2.8	–
Euler- B , $\Delta=\tau g$	0.21	0.68	1.1	1.4	3.1	0.93
bilinear, $\Delta=\tau g$	9.5×10^{-3}	0.48	0.94	1.3	2.8	2.25

HAR data (acquisition command and archive hash are given) and not the third-party TIDES code, which is obtained at the pinned commit under its own license.

I. Additional toy tables

Tables 17 to 22 list the toy measurements referenced in Section 6: one fix at a time, splitting, the two-layer cascade, the stiff mode, new observations with the exact decomposition, and readouts under a density change.

Table 19. Two-layer linear cascade $\dot{h}_1 = -ah_1 + bu$, $\dot{h}_2 = -ch_2 + dh_1$ under interval splitting (P1-COMP-01; $b=d=1$; held input of FIRST_RUN unit 0). Columns m : error of the layerwise computation (layer 2 holds the layer-1 endpoint value) against the exact coupled solution, $\|h_2^{\text{lw}} - h_2\|/\|h_2\|$ over base times. “change” is the layerwise change $m=8$ vs. $m=1$; the last column is the largest change of the exact coupled solution over m . No external observation is added in this experiment.

case	$m=1$	$m=2$	$m=4$	$m=8$	slope	change	exact inv.
$a=1, c=0.5$	0.19	0.10	0.052	0.026	-0.95	0.15	1.7×10^{-15}
$a=c=1$	0.20	0.11	0.055	0.028	-0.93	0.16	7.1×10^{-16}

Table 20. Stiff scalar mode (FR-E4-N2; $g = \ln 2$, base step 0.5, held sine input). Bilinear is stable ($|\bar{A}| < 1$) for every step, but for $z = \Delta\lambda < -2$ its transition is negative (sign-alternating transients) and tends to -1 , whereas the exact decay e^z tends to 0. The last two columns are the output artifact against exact ZOH. The stable/sign/decay columns are an exploratory re-reading of logged values.

λ	m	z	$\bar{A}_{\text{bil}}$	e^z	stable	sign-alt.	bilinear	Euler- B
-5	1	-1.73	+0.072	0.18	yes	no	0.098	1.1
-5	4	-0.43	+0.644	0.65	yes	no	4.7×10^{-3}	0.23
-5	16	-0.11	+0.897	0.90	yes	no	2.9×10^{-4}	0.055
-5	32	-0.05	+0.947	0.95	yes	no	7.3×10^{-5}	0.027
-50	1	-17.33	-0.793	3.0×10^{-8}	yes	yes	0.54	16
-50	4	-4.33	-0.368	0.013	yes	yes	0.015	3.4
-50	16	-1.08	+0.297	0.34	yes	no	2.1×10^{-8}	0.64
-50	32	-0.54	+0.574	0.58	yes	no	8.5×10^{-9}	0.30
-500	1	-173.29	-0.977	5.5×10^{-76}	yes	yes	1.1	170
-500	4	-43.32	-0.912	1.5×10^{-19}	yes	yes	0.62	42
-500	16	-10.83	-0.688	2.0×10^{-5}	yes	yes	2.0×10^{-3}	9.8
-500	32	-5.42	-0.461	4.4×10^{-3}	yes	yes	1.3×10^{-11}	4.4

dt	exact ZOH	Euler- B	bilinear
0.01	8.0×10^{-15}	3.5×10^{-3}	0
0.5	0	0.18	1.1×10^{-16}
8	0	4.6	5.6×10^{-6}

Lower block (FR-E4-N3): relative steady-state error under a held unit input, $\lambda = -1$.

Table 21. Adding new observations of the same signal (FR-E2-NEWOBS; toy, 32 units). The change $T = y_m - y_1$ at base times splits *exactly* as $T = P + R$, where P is the change of the continuous-time solution on the new held path (the change in the held-input reconstruction; no task or Bayesian information gain is implied) and R is the change of the discretization artifact. Hence $\|T\|^2 = \|P\|^2 + \|R\|^2 + 2\langle P, R \rangle$; we report these three terms as fractions of $\sum_{\text{units}} \|T\|^2$ (they sum to 1) and the per-unit range of the cross fraction. P uses the exact-ZOH output as a proxy for the continuous-time held-path solution (validated against RK4 to the bound shown). Median total change is relative to $\|y_1\|$. EXPLORATORY re-aggregation; the cross term shows that no per-unit “share” is well defined, and none of these fractions is a causal contribution.

variant	m	median $\ T\ /\ y_1\ $	$\ P\ ^2$	$\ R\ ^2$	$2\langle P, R \rangle$	cross range
exact ZOH, $\Delta=dt\ g$	2	0.11	artifact $\leq 1.7 \times 10^{-13}$ vs. RK4 (proxy reference)			
exact ZOH, $\Delta=dt\ g$	8	0.19	artifact $\leq 1.7 \times 10^{-13}$ vs. RK4 (proxy reference)			
Euler- B , $\Delta=dt\ g$	2	0.15	0.173	0.567	+0.260	$[-0.93, +0.42]$
Euler- B , $\Delta=dt\ g$	8	0.25	0.198	0.533	+0.269	$[-1.37, +0.42]$
bilinear, $\Delta=dt\ g$	2	0.11	0.849	0.025	+0.126	$[-1.04, +0.34]$
bilinear, $\Delta=dt\ g$	8	0.19	0.886	0.013	+0.101	$[-0.69, +0.29]$
exact ZOH, $\Delta=\tau g$	2	0.42	0.071	1.260	-0.331	$[-1.06, +0.34]$
exact ZOH, $\Delta=\tau g$	8	1.1	0.025	1.223	-0.248	$[-1.02, +0.03]$
Euler- B , $\Delta=\tau g$	2	0.42	0.053	1.225	-0.278	$[-0.72, +0.32]$
Euler- B , $\Delta=\tau g$	8	1.1	0.020	1.202	-0.222	$[-0.75, +0.04]$
bilinear, $\Delta=\tau g$	2	0.42	0.072	1.260	-0.331	$[-1.10, +0.35]$
bilinear, $\Delta=\tau g$	8	1.1	0.025	1.222	-0.247	$[-1.04, +0.03]$

Table 22. Readouts under a density change (FR-E3-READOUT; toy, 32 units). Columns 2–3: scale-normalized change $|r_8 - r_1|/s$ when only the first half of the base intervals is split into 8 held sub-steps (physical path unchanged); s is the RMS base-time output at $m=1$ ($\times K_1$ for the sample sum). Column 4: scale-normalized error to the continuous-time truth when new observations are added to the first half only, $m=1 \rightarrow 8$. Medians over units. The exact held-path integral is defined only for exact-ZOH states.

readout	split, exact ZOH $\Delta=dt\,g$	split, Euler- B $\Delta=\tau g$	new obs., exact ZOH
last sample	8.9×10^{-17}	0.19	0.085→0.099
sample sum	1.5	2.9	–
sample mean	0.30	0.46	0.060→0.28
right-endpoint $\sum y_k\,dt_k$	0.026	0.24	0.060→0.043
trapezoidal endpoint sum	0.011	0.25	0.039→0.023
exact held-path integral	1.5×10^{-16}	n/a	0.039→0.032